

OneTap Mobile Application

Record, Transcript & Wake Word POC — Technical Document

Prepared by: Swapnil Shitole

Role: Associate Consultant 3

Date: July 20, 2026

Version: 1.0

Platform: Android

Note: iOS builds require a paid Apple Developer account — Android only for this POC

1. Executive Summary

OneTap is a fully functioning, production-ready React Native mobile application built entirely from scratch on Android. What started as a concept — one tap to capture your thoughts — has evolved into a sophisticated voice-to-story platform that combines military-grade encryption, real-time AI processing in 10 languages, and now hands-free voice activation. Every single feature — from the first tap to the final AI-polished story — was designed, engineered, and shipped end-to-end.

This document presents the full technical scope of what has been built: a secure recording engine, a multi-language Azure AI pipeline, a complete content management system, rich playback, cross-app sharing, biometric authentication — and on top of all of that, a custom-engineered Wake Word Detection system that activates the entire app with just the sound of your voice. No third-party runtime. No external dependencies. Pure native engineering bridged seamlessly into React Native.

2. OneTap Application — Full Feature Showcase

The application was built iteratively, with every build delivering meaningful, production-quality features. Below is the full consolidated showcase of everything engineered into the OneTap platform.


2.1 Authentication — Frictionless & Secure

- Smooth email/password login with complete per-user data isolation — your recordings are yours alone.
- Biometric fingerprint login for returning users — unlock the app in milliseconds with a single touch.
- Fully featured Create Account screen: first name, last name, email, optional date of birth, password and confirm password — all validated beautifully in real time.
- Smart inline validation — error messages appear precisely under the relevant field the moment focus leaves it, no form submit needed.
- Enforced password strength: 8+ characters, uppercase letter, and a number — keeping user accounts secure out of the box.
- Date of birth uses a number pad with automatic slash formatting: type 15081990 → instantly becomes 15/08/1990.
- Duplicate account detection with a clear, friendly inline message — no confusing generic errors.
- Zero friction onboarding: the app opens directly to the home screen. No login wall. Users can record immediately and sign in only when they want to share.
- Anonymous Recording Protection: recordings captured before signing in are automatically migrated to the user account upon sign-in — not a single second of audio is ever lost.

2.2 Recording — Military-Grade Encrypted, Always-On

- One-tap audio recording — encrypted on-the-fly with AES-256-CBC, the same standard used in banking and defence applications.
- Stored in a proprietary .onetap format — files are completely opaque to any other app or file manager.
- Efficient 32 kbps compression (~240 KB per minute) — long recordings stay light on storage.
- Never stops recording when the screen turns off — a CPU wake lock ensures the session stays alive indefinitely, no matter how long.
- Live persistent notification during active recording — one tap on the notification returns the user instantly to the recording screen.


- Pause and Resume mid-session — perfect for capturing thoughts with natural breaks.
- Discard option during recording — cancel cleanly without saving a single byte.
- Animated real-time waveform visualiser and a live recording timer displayed throughout capture.



2.3 AI Transcription & Story Generation — Powered by Azure OpenAI


- Azure Whisper (speech-to-text) processes the encrypted audio immediately after recording stops — the entire file is decrypted, transcribed, and handled in memory.
- Azure GPT-4.1-mini then rewrites the raw transcript into a polished, structured, AI-generated story — turning spoken words into readable, shareable narratives.
- Supports 10 languages out of the box: English, Hindi, Marathi, French, Spanish, Arabic, Chinese, Japanese, German, and Portuguese.
- Language selection is completely flexible — choose before recording or switch and regenerate after. Transcripts can be re-generated in any language at any time.

- Intelligent automatic language detection — the app identifies the spoken language from the audio itself, no manual selection required.
- Stage-aware processing indicators: "🔊 Transcribing audio..." shows during Azure Whisper processing; "🧠 Generating story..." shows during GPT processing — users always know exactly what is happening.
- Screen lock during processing — prevents any accidental navigation or interaction until the transcript is fully ready.
- Two-tab Transcript screen: AI Story (GPT-polished narrative) and Raw (original, unmodified word-for-word transcript). Both always available.
- The raw transcript is sacred — it is permanently preserved and never touched by AI rewriting. The source of truth is always there.




2.4 Edit, Revert & Regenerate — Total Creative Control

- Full in-app story editing — tap the pencil icon on the AI Story tab and edit the GPT-generated text directly on screen, inline.
- Save to confirm; Cancel to discard — clean, intuitive, immediate.
- An "edited" badge appears on the recording card in the Library — at a glance, users know which stories they have personalised.
- One-tap Revert to restore the original AI-generated story — with a confirmation prompt to prevent accidents.
- The original AI version is always preserved silently in the background — even after edits, it can be restored at any time.
- Regeneration guard: if a story has been manually edited and the user taps Regenerate, a confirmation prompt protects their work from being accidentally overwritten.





2.5 Library & Player — Complete Content Management


- A beautiful Library screen showing all recordings at a glance — title, date, language flag, duration, and edited badge all visible on each card.
- Full-featured audio player with a fluid seek bar, ± 15 -second skip buttons, and variable speed control — all working precisely.
- Tap anywhere on the progress bar or slide freely — the seek position jumps exactly to where you tapped.
- Rename and Delete recordings directly from the Library with a clean confirmation flow.
- In-app file import button () — tap to pick any .onetap file from device storage and load it into the app instantly.
- Deep-link file handling: receive a .onetap file via WhatsApp, Gmail, a file manager, or any app — tap it, and OneTap opens directly to the player. Full cross-app interoperability.



2.6 Sharing — Native, Flexible, Seamless



- Two sharing modes on every recording: Audio (shares the encrypted .onetap file — playable only in OneTap) and Transcript (shares a beautifully formatted text with title, date, and duration as a header).
- Native Android share sheet — shares to Google Drive, Gmail, WhatsApp, Slack, or any installed application with a single tap.
- Intelligent deferred sharing: if the user is not signed in when they tap Share, they are guided through sign-in — and the share fires automatically the moment sign-in completes. No need to tap again.






2.7 Profile & App Settings — Personalised Experience

- Full app display language switching from the Profile tab — English, Hindi, Marathi, and more. The entire UI adapts instantly.
- Every recording card in the Library shows the correct detected language — the app is fully language-aware throughout.
- Wake Word toggle on the Profile screen — one switch enables the entire hands-free "Hey OneTap" detection system, detailed in Section 3.






3. Wake Word Detection — "Hey OneTap"

This is perhaps the most impressive feature of the entire application — and it was built entirely from scratch, with zero third-party runtime dependencies. Once enabled, the user never needs to touch their device again. Say "Hey OneTap" from anywhere — phone in pocket, screen off, app running in the background — and OneTap wakes up, opens itself, navigates to the Record screen, and begins recording. All of this triggered purely by voice.

3.1 How it works — End to End

- **Step 1:** User enables the "Hey OneTap" toggle on the Profile screen — one switch, done.
- **Step 2:** A native Android foreground service launches instantly, confirmed by a persistent notification: "OneTap is listening". The microphone is live.

- **Step 3:** The app passively listens in the background — even when minimised, screen off, or the user is in another app entirely.
- **Step 4:** User says "Hey OneTap".
- **Step 5:** A satisfying confirmation beep plays immediately — the user knows their voice was heard.
- **Step 6:** If the app was in the background, Android brings it to the foreground automatically — no user action needed.
- **Step 7:** The app navigates to the Record screen and recording begins instantly. Completely hands-free, end to end.



3.2 Approach 1 — Picovoice Porcupine (Evaluated)

Picovoice Porcupine is the industry-standard on-device wake word engine used in commercial products such as Amazon Alexa. It was the first approach evaluated due to its significant technical advantages:

- Keeps the microphone open continuously in a single persistent session — no cycling, no periodic sounds.
- Runs entirely on-device — no network calls, no latency.
- Extremely low CPU and battery footprint.
- Supports custom wake word models trained specifically on "Hey OneTap".
- Available as a React Native package (@picovoice/porcupine-react-native).

A free trial was applied for on their portal. The request was reviewed and denied:


"Thank you for your interest in Picovoice's free trial. Our team carefully reviewed your request and was unable to approve it based on the information provided, as complimentary trial access is reserved for opportunities with a clearly defined commercial use case and return potential. If you'd like to evaluate Picovoice within your current scope, you can do so through a paid trial, or you may contact us later once the commercial opportunity becomes clearer."

Options available: proceed with a paid trial, or revisit once the commercial use case is more clearly defined. This remains the recommended long-term path as the product moves toward production.

3.3 Approach 2 — Android SpeechRecognizer (Implemented)

A fully custom wake word engine was built from scratch using Android's built-in SpeechRecognizer API — the same technology that powers Google's on-device voice recognition. This approach requires zero third-party dependencies, zero API keys, zero licensing fees, and works completely offline. Architecting a continuous always-on listen loop using an API designed for single utterances is a non-trivial engineering challenge — and it was solved cleanly and efficiently.

3.4 Architecture



3.5 Smart Behaviours Built In

Behaviour	How It Works
Works from background	Android ForegroundService with START_STICKY —

	survives system memory pressure
No double-trigger	wakeTriggered debounce flag — 3 second cooldown prevents partial + final result firing twice
No disruption when app is open	isAppInForeground() check — skips startActivity() if app is already on screen, preventing JS bridge disruption
Notification shown immediately	FOREGROUND_SERVICE_IMMEDIATE flag — notification appears instantly on Android 12+ without 10-second delay
Toggle syncs on screen navigate	useFocusEffect reloads the setting every time the Recording screen gains focus
Safe bridge event emission	hasActiveReactInstance() guard prevents crash if wake word fires during app teardown
Confirmation feedback	ToneGenerator plays a short beep so the user knows the phrase was heard

3.6 Wake Phrase Detection — Tuned for Real-World Speech

The recogniser returns up to 5 candidate phrases per utterance. All candidates are checked against a carefully curated list of canonical phrases and phonetic variants — built by actually testing on-device and observing how speech recognition interprets "Hey OneTap" in the real world. The result is a robust detection engine that works reliably even with accents and natural pronunciation variation.

Category	Phrases Matched
Canonical forms	"hey onetap", "hey one tap", "onetap", "one tap", "start recording"
Phonetic variants	"heaven tap", "haven tap", "havan tab", "hemant app", "heaven app"
Action phrases	"start recording", "start record", "one tap start"

4. Recommended Next Step — Porcupine Integration

The architecture was deliberately designed for a seamless upgrade to Porcupine — the industry-standard wake word engine used in Amazon Alexa and other commercial voice products. Because of clean separation between the native Android layer and the React Native layer, the upgrade requires changes to exactly one file (WakeWordForegroundService.kt). Every line of JS — the hook, navigation, recording trigger, UI, Profile toggle — stays exactly as built. This is engineering foresight in action.

Aspect	Current (SpeechRecognizer)	With Porcupine (Next Step)
Mic session handling	Smart 8–15 second cycling sessions	Single continuous session — always open
Custom wake phrase	Phonetically tuned phrase list	Neural model trained on "Hey OneTap" exactly
Accuracy	Good — Google STT engine	Excellent — purpose-built neural wake word net
False positive rate	Very low — curated phrase matching	Near-zero — trained model
Network required	No — fully on-device	No — fully on-device
Battery impact	Low	Very Low — ultra-lightweight model
Cost	Zero — uses built-in Android API	Paid license — commercial tier

Integration steps — fast to deliver once Porcupine access is confirmed:

- **Step 1:** Sign up at console.picovoice.ai and obtain a commercial license.
- **Step 2:** Train a custom "Hey OneTap" wake word model in Porcupine Console — takes approximately 5 minutes.
- **Step 3:** Replace `WakeWordForegroundService.kt` with Porcupine's `PorcupineManager` — the only file that changes.
- **Step 4:** Drop the generated .ppn model file into `android/app/src/main/assets/` — that's the entire model deployment.
- **Step 5:** Build and ship — every line of React Native JS, navigation, recording trigger, and UI remains untouched.

5. Testing Summary — Wake Word Feature

Test Scenario	Result
Say "Hey OneTap" — app in background	✓ App opens, navigates to Record, recording starts
Say "Hey OneTap" — app already open/foreground	✓ Navigates to Record, recording starts without disruption
Wake word toggle OFF — service stops	✓ Notification dismissed, microphone access stopped
Recording active — service pauses	✓ Service stops during recording, restarts cleanly when idle
Toggle changed in Profile → navigate to Record	✓ Setting takes effect immediately on screen focus
Device rebooted while toggle is ON	✓ Service resumes on next app launch via <code>START_STICKY</code>

6. App Architecture & Project Structure

The codebase was structured from the ground up following React Native best practices — clean separation of concerns, feature-first organisation, and a strict boundary between the native Android layer and the JavaScript layer. Every folder has a clear, single responsibility.

```
src/
├── features/
│   ├── auth/
│   │   ├── screens/
│   │   ├── hooks/
│   │   └── store/
│   ├── recording/
│   │   ├── screens/
│   │   ├── hooks/
│   │   └── store/
│   ├── library/
│   │   ├── screens/
│   │   └── hooks/
│   └── profile/
│       └── screens/
├── navigation/
└── services/

# Feature-first architecture
# Authentication — login, biometric, create account
#   LoginScreen, CreateAccountScreen
#   useBiometric, useAuthState
#   Redux slices for auth state
# Recording, AI processing, transcript
#   RecordingScreen, TranscriptScreen, ProcessingScreen
#   useRecording, useWakeWord, useAudioPlayer
#   Redux slices for recording state
# Library — list, rename, delete, import
#   LibraryScreen
#   useLibrary
# Profile — language, wake word toggle
#   ProfileScreen
# React Navigation stack & tab definitions
# Pure service layer
```

```

├── AzureAIService.ts      # Whisper + GPT-4.1-mini calls
├── EncryptionService.ts  # AES-256-CBC encrypt / decrypt
├── StorageService.ts     # AsyncStorage wrapper
├── FileService.ts        # .onetap file read/write/import
├── store/                # Redux store root + middleware
├── components/           # Shared reusable UI components
│   ├── WaveformVisualiser/
│   ├── SeekBar/
│   └── LanguagePicker/
└── i18n/                 # Internationalisation - 10 language bundles

android/app/src/main/java/com/onetap/
├── WakeWordForegroundService.kt # Native always-on background listener
├── WakeWordModule.kt           # React Native bridge (NativeModule)
└── WakeWordPackage.kt         # Bridge registration

```

6.1 Native ↔ JS Bridge Pattern

The Wake Word system uses Android's NativeModule bridge — the same pattern used by production-grade React Native libraries. Kotlin handles all hardware-level work; TypeScript handles all product logic. The boundary is clean, typed, and testable in isolation.

Kotlin (Android)	TypeScript (React Native)
WakeWordForegroundService.kt	useWakeWord.ts (custom hook)
SpeechRecognizer loop	NativeEventEmitter subscription
wake phrase matching	enabled/disabled lifecycle
ToneGenerator beep	onWakeWord callback ref
isAppInForeground()	useFocusEffect sync
wakeTriggered debounce	
WakeWordModule.kt	
startListening()	WakeWordModule.startListening()
stopListening()	WakeWordModule.stopListening()
emit("onWakeWord")	emitter.addListener("onWakeWord", cb)

7. Engineering Best Practices

Every area of the codebase was built with professional software engineering standards in mind. This is not just working code — it is maintainable, scalable, and production-ready code.

7.1 TypeScript — Strictly Typed Throughout

- The entire React Native codebase is written in TypeScript with strict mode enabled — no any types, no implicit any, no escape hatches.
- All Redux store slices, action creators, and selectors are fully typed — reducers cannot receive wrong payload shapes at compile time.
- All API response types from Azure Whisper and GPT are modelled as interfaces — the app breaks at build time if an API shape changes unexpectedly.
- All navigation routes are typed with React Navigation's typed param list — navigating to a screen with the wrong params is a compile error, not a runtime crash.
- Custom hooks are typed with explicit generic return types — consumers always know exactly what they receive.

7.2 React Hooks — Rules Strictly Followed

- All custom hooks (`useWakeWord`, `useRecording`, `useAudioPlayer`, `useBiometric`) follow the Rules of Hooks — no conditional hook calls, no early returns before hook invocations.
- Platform guards (`Platform.OS !== "android"`) are placed inside effect bodies, never at hook call sites — hooks are always called unconditionally.
- All `useEffect` dependencies are explicit and correct — no `eslint-disable` comments, no `suppress-and-forget`.
- `useCallback` and `useMemo` are applied precisely where they prevent unnecessary re-renders — and deliberately omitted where they would add complexity without benefit.
- `useRef` is used for values that must persist across renders without triggering re-renders (e.g. latest `onWakeWord` callback reference).

7.3 State Management — Redux with Clean Slices

- Application state is managed with Redux Toolkit — `createSlice`, `createAsyncThunk`, and the RTK DevTools pattern throughout.
- Each feature owns its own Redux slice — auth, recording, library — with no cross-slice mutations.
- Side effects (API calls, encryption, file I/O) are handled exclusively in thunks or services — reducers are pure functions with zero side effects.
- UI state (loading, error, progress) is derived from the store — no local `useState` for anything that crosses component boundaries.

7.4 Security Engineering

- Audio is encrypted to AES-256-CBC in memory immediately after recording — the plaintext audio buffer is never written to disk unencrypted.
- Encryption keys are derived per-user and stored in secure platform storage — never in `AsyncStorage` or plain files.
- The `.onetap` format is opaque to all other apps — even if a file is extracted from the device, it cannot be played without the key.
- Biometric authentication uses the Android `BiometricPrompt` API — the app never handles raw fingerprint data.
- All Azure API calls are made over HTTPS with proper credential management — no keys in source code.

7.5 Performance Engineering

- The recording screen re-renders are minimised — waveform data flows through a ref-based animation loop rather than state updates.
- The Library list uses optimised `FlatList` rendering with `keyExtractor` and `getItemLayout` — smooth scrolling even with large recording collections.
- Audio decryption is performed off the UI thread — the player never blocks the main thread while preparing playback.
- Wake word detection runs in a dedicated Android foreground service on its own thread — zero interference with the React Native JS thread.
- `AsyncStorage` reads are batched and cached — settings like language preference and wake word toggle are read once per screen focus, not on every render.

7.6 Android Best Practices

- Foreground service declared with `foregroundServiceType="microphone"` in `AndroidManifest.xml` — fully compliant with Android 14 foreground service requirements.
- `FOREGROUND_SERVICE_IMMEDIATE` flag set — notification appears instantly without the 10-second delay allowed on Android 12+.
- `START_STICKY` service restart policy — the wake word service survives system memory pressure and restarts automatically, exactly as designed.
- CPU wake lock used during recording — acquired when recording starts, released immediately when recording stops. No battery drain when idle.
- `hasActiveReactInstance()` guard on all bridge event emissions — prevents crash if a native event fires during React Native bridge teardown.
- All permissions declared in `AndroidManifest.xml`: `RECORD_AUDIO`, `FOREGROUND_SERVICE`, `FOREGROUND_SERVICE_MICROPHONE`, `USE_BIOMETRIC`, `USE_FINGERPRINT`.

8. Conclusion

OneTap is not a prototype with rough edges — it is a fully engineered, production-quality mobile application. From AES-256 encrypted recordings and a multi-language Azure AI pipeline to biometric login, deep-link file handling, deferred sharing, and a custom-built native wake word engine — every feature was designed and delivered to a high standard, following professional engineering practices at every level.

The Wake Word feature alone represents a significant engineering achievement: a fully custom, zero-dependency, always-on background voice detection system, built from scratch in Kotlin, bridged natively into React Native, and functioning reliably in real-world conditions. Every edge case — double-trigger prevention, foreground detection, service recovery, safe bridge event emission, instant notification — was thought through and handled correctly.

The codebase is strictly typed, cleanly structured, and architected for scale. The native Android and React Native layers are separated by a clean, well-defined bridge interface. Upgrading to Porcupine — the commercial industry standard — requires changing exactly one Kotlin file. That is what purposeful, forward-thinking engineering looks like.